

Model Setup

Consider a standard OLS regression applied to the dataset $d_k = \{y_{i,k}, x_{i1,k}, \dots, x_{ip,k}\}_{i=1}^n$ in the test of anomaly k

$$y_{i,k} = \beta_{0,k} + \beta_{1,k}x_{i1,k} + \dots + \beta_{p,k}x_{ip,k} + \varepsilon_{i,k} \quad (1)$$

with independent residual errors $\varepsilon_{i,k} \sim \mathcal{N}(0, \sigma^2)$ ¹. In this setup, $\beta_{0,k}$ is the coefficient of interest. Let $\mathcal{M}_{\text{null},k}$ denote the null model that restricts $\beta_{0,k} \leq 0$

$$\mathcal{M}_{\text{null},k} := \{y_{i,k} = \beta_{0,k} + \beta_{1,k}x_{i1,k} + \dots + \beta_{p,k}x_{ip,k} + \varepsilon_{i,k}, \quad \beta_{0,k} \leq 0\}$$

and $\mathcal{M}_{\text{alt},k}$ denote the alternative model that restricts $\beta_{0,k} > 0$

$$\mathcal{M}_{\text{alt},k} := \{y_{i,k} = \beta_{0,k} + \beta_{1,k}x_{i1,k} + \dots + \beta_{p,k}x_{ip,k} + \varepsilon_{i,k}, \quad \beta_{0,k} > 0\}.$$

Our goal is to compute the posterior probability of the null model, $\Pr(\mathcal{M}_{\text{null},k} | d_k)$, given the data d_k .

We place the non-informative (Jeffreys) prior on the other coefficients $\boldsymbol{\beta}_{1:p,k} = (\beta_{1,k}, \dots, \beta_{p,k})^\top$ and the error variance σ^2

$$\pi(\boldsymbol{\beta}_{1:p,k}, \sigma^2) \propto \frac{1}{\sigma^2}.$$

For $\beta_{0,k}$, we use a mixed one-sided prior derived from the base normal distribution $\mathcal{N}(0, \tau^2)$, where τ is the prior scale parameter, that encodes a belief about its sign

$$\pi(\beta_{0,k}) = \pi_{\text{null}} 2\tau^{-1} \phi(\beta_{0,k} / \tau) \mathbb{I}(\beta_{0,k} \leq 0) + (1 - \pi_{\text{null}}) 2\tau^{-1} \phi(\beta_{0,k} / \tau) \mathbb{I}(\beta_{0,k} > 0),$$

where $\phi(\cdot)$ is the standard normal distribution² and $\mathbb{I}(\cdot)$ is an indicator function equal to one if the condition is satisfied and zero otherwise. The hyperparameter $0 \leq \pi_{\text{null}} \leq 1$ is the prior probability of $\mathcal{M}_{\text{null},k}$, and $\tau > 0$ is a scale parameter for $\beta_{0,k}$. This prior is essentially a normal distribution $\mathcal{N}(0, \tau^2)$ for $\beta_{0,k}$ truncated to negative or positive values, with a mixing weight π_{null} favoring the negative side and $1 - \pi_{\text{null}}$ the positive side.

¹The assumption of independent model errors is not essential for our framework as models with more sophisticated assumptions on data dependence fit into this framework straightforwardly.

²Thus, $2\tau^{-1}\phi(\beta_{0,k} / \tau)$ is a half-normal distribution on either side of zero.

Posterior Inference Approach

The full likelihood function for the model (1) conditioned on the data d_k is

$$L(\beta_{0,k}, \boldsymbol{\beta}_{1:p,k}, \sigma^2 | d_k) = (2\pi\sigma^2)^{-n/2} \exp\left[-\frac{1}{2\sigma^2}(\mathbf{y} - \mathbf{X}[\beta_{0,k}, \boldsymbol{\beta}_{1:p,k}])^\top (\mathbf{y} - \mathbf{X}[\beta_{0,k}, \boldsymbol{\beta}_{1:p,k}])\right] \quad (2)$$

where $\mathbf{y} = \{y_{1,k}, \dots, y_{n,k}\}^\top$, $\mathbf{X} = [\mathbf{1}, \mathbf{Z}]$ with $\mathbf{1}$ being the column of ones (to accommodate $\beta_{0,k}$) and \mathbf{Z} containing the p covariates. Integrate the likelihood function (2) with respect to $\boldsymbol{\beta}_{1:p,k}$ and σ^2 under the Jeffreys prior to obtain the likelihood function for $\beta_{0,k}$

$$L(\beta_{0,k} | d_k) = c[1 + t^2(\beta_{0,k}) / \nu]^{-\frac{\nu+1}{2}}$$

where

$$t(\beta_{0,k}) = \frac{\beta_{0,k} - \hat{\beta}_{0,k}}{\text{SE}(\hat{\beta}_{0,k})}, \quad \nu = n - p - 1, \quad c = \frac{\Gamma(\frac{\nu+1}{2})}{\Gamma(\frac{\nu}{2})\sqrt{\pi\nu}\text{SE}(\hat{\beta}_{0,k})}(2\pi)^{-\nu/2}|\mathbf{Z}^\top \mathbf{Z}|^{-1/2}.$$

Here, $t(\beta_{0,k})$ is the usual t -statistic and $\text{SE}(\hat{\beta}_{0,k})$ is the standard error of the OLS estimate $\hat{\beta}_{0,k}$. Consequently, $L(\beta_{0,k} | d_k)$ is (up to the constant c) the Student t -distribution for $\beta_{0,k}$ with ν degrees of freedom, centered at $\hat{\beta}_{0,k}$ and having scale $\text{SE}(\hat{\beta}_{0,k})$.

Then, we compute the marginal likelihood under each sub-model, $L_{\text{null}}(d_k)$ and $L_{\text{alt}}(d_k)$, by integrating the conditional likelihood $L(\beta_{0,k} | d_k)$ over the half-normal prior on the restricted domain of $\beta_{0,k}$ for each model, respectively,

$$L_{\text{null}}(d_k) = \int_{-\infty}^0 L(\beta_{0,k} | d_k) 2\tau^{-1} \phi(\beta_{0,k} / \tau) d\beta_{0,k}$$

and

$$L_{\text{alt}}(d_k) = \int_0^{\infty} L(\beta_{0,k} | d_k) 2\tau^{-1} \phi(\beta_{0,k} / \tau) d\beta_{0,k}.$$

The posterior probability of $\mathcal{M}_{\text{null},k}$ is

$$\Pr(\mathcal{M}_{\text{null},k} | d_k) = \frac{\pi_{\text{null}} L_{\text{null}}(d_k)}{\pi_{\text{null}} L_{\text{null}}(d_k) + (1 - \pi_{\text{null}}) L_{\text{alt}}(d_k)}.$$